

On using a DSL Approach Performance Portability of the LFRic Weather and Climate Model

Dr Chris Maynard


Scientific Software Engineer - Met Office

Associate Professor of Computer Science – University
of Reading

www.aces.cs.reading.uk

Paper DOI

<https://doi.org/10.1016/j.jpdc.2019.02.007>



PSyclone

MS02: Kavcic: Wed 1300 –

PSyclone and its Use in LFRic

MS29: Ford: Thur 1515 – *PSyIR: the PSy Intermediate Representation*

LFRic

MS23: CMM: Thurs 1245 - *Scalable Linear Solvers for Next Generation Weather and Climate Models*

CSM07: Poster : *Building a Performance Portable Software System for the Met Office's Weather and Climate Model, LFRic*

Fortran – high level language

Abstraction of the numerical mathematics

Implementation and architecture is hidden

Code – text which conforms to the semantics and syntax of the language definition

Compiler transforms code into

real(kind=r_def), dimension(nqp_h), intent(in) :: wqp_h
real(kind=r_def), dimension(nqp_v), intent(in) :: wqp_v

! Internal variables

integer :: df, df2, k, ik
integer :: qp1, qp2

real(kind=r_def), dimension(ndf_chi) :: chi1_e, chi2_e, chi3_e
real(kind=r_def) :: integrand

real(kind=r_def), dimension(ndf_chi) :: dj

real(kind=r_def), dimension(3,3,nqp_h,nqp_v) :: jac

! loop over layers: Start from 1 as in this loop k is not an offset

do k = 1, nlayers
 ik = k + (cell-1)*nlayers


! indirect the chi coord field here

do df = 1, ndf_chi
 chi1_e(df) = chi1(map_chi(df) + k - 1)
 chi2_e(df) = chi2(map_chi(df) + k - 1)
 chi3_e(df) = chi3(map_chi(df) + k - 1)
 end do

call coordinate_jacobian(ndf_chi, nqp_h, nqp_v, chi1_e, chi2_e, chi3_e, &
 diff_basis_chi, jac_dij)

Separation of concerns

Abstraction is *broken* by parallel/performance/memory features exposed
Hacked back together with
MPI, OMP, Open ACC, OpenCL, CUDA, PGAS, SIMD, compiler directives
Libraries, languages (exts), directives and compiler (specific) directives



Scientific programming


Find numerical solution (and estimate of the uncertainty) to a (set of) mathematical equations which describe the action of a physical system

Parallel programming and optimisation are the methods by which large problems can be solved faster than real-time.




SEPARATION OF CONCERN

Don't let your plumbing code pollute your software.



Cubed Sphere → no singular poles lon-lat
Unstructured mesh → can use other meshes
Mixed finite element scheme – C-Grid
Exterior calculus *mimetic* properties
Semi-implicit in time



Alg layer – high level expression of operations on global fields

Kernel layer – low level Explicit operation on a single column of data


Code has to follow set of rules
(PSyKAI API is DSL)

Parallelisation System

Horizontal looping and parallel code.

Can generate parallel code according to rules

- PSyKAI



```
call invoke(
    held_suarez_kernel_type(
        rhs_heoldsuarez(igh_u) ,
        rhs_heoldsuarez(igh_t) ,
        state_n(igh_u) ,
        state_n(igh_t) ,
        state_n(igh_d) ,
        chi, qr) ,
    enforce_bc_kernel_type(
        rhs_heoldsuarez(igh_u) )
)
```

invoke() Do this in parallel kernels single column operations fields data parallel global fields

Multiple kernels in single invoke → scope of ordering/parallel communication, etc

Embed metadata as (compilable) Fortran, but it doesn't get executed

Data Access descriptors

Explicitly describe kernel arguments

Richer information than Fortran itself

```
!> The type declaration for the kernel. Contains the metadata needed by the Psy layer
type, public, extends(kernel_type) :: exner_gradient_kernel_type
  private
    type(arg_type) :: meta_args(3) = (
      arg_type(GH_FIELD, GH_INC, W2),
      arg_type(GH_FIELD, GH_READ, W3),
      arg_type(GH_FIELD, GH_READ, ANY_SPACE_9)
    )
    type(func_type) :: meta_funcs(3) = (
      func_type(W2, GH_BASIS, GH_DIFF_BASIS),
      func_type(W3, GH_BASIS),
      func_type(ANY_SPACE_9, GH_BASIS, GH_DIFF_BASIS)
    )
    integer :: iterates_over = CELLS
    integer :: gh_shape = GH_QUADRATURE_XYoz
    ! gh_shape replaces evaluator_shape
    integer :: evaluator_shape = QUADRATURE_XYoz
contains
  procedure, nopass ::exner_gradient_code
end type
```

Python code generator

Parser, transformations, generation

Controls parallel code (MPI/OpenMP and OpenACC)

Potentially other programming models

e.g. OpenCL for FPGA

MS02: Kavcic: Wed 1300 –

PSyclone and its Use in LFRic

MS29: Ford: Thur 1515 – *PSyIR: the PSy*

Intermediate Representation

Update halos
YAXT→MPI

colouring from
infrastructure

OpenMP
workshare across
cells in colour

```
!
IF (chi_proxy(3)%is_dirty(depth=1)) THEN
  CALL chi_proxy(3)%halo_exchange(depth=1)
END IF
!
CALL rhs_heldsuarez_proxy%vspace%get_colours(ncolour, ncp_colour, cmap)
!
DO colour=1,ncolour
  !$omp parallel default(shared), private(cell)
  !$omp do schedule(static)
  DO cell=1,ncp_colour(colour)
    !
    call held_suarez_code(nlayers, rhs_heldsuarez_proxy%data, &
      ! ...
      )
  END DO
  !$omp end do
  !$omp end parallel
END DO
!
! Set halos dirty for fields modified in the above loop
!
CALL rhs_heldsuarez_proxy%set_dirty()
```

kernel call for single
column. Args are
arrays and scalars

Single kernel invoke

```
Transforming invoke 'invoke_26_rtheta_kernel_type' ...
Schedule[invoke='invoke_26_rtheta_kernel_type' dm=False]
    Loop[type='',field_space='w0',it_space='cells', upper_bound='ncells']
        KernCall rtheta_code(rtheta,theta,wind) [module_inline=False]
```

Apply distributed memory

```
Transforming invoke 'invoke_26_rtheta_kernel_type' ...
Schedule[invoke='invoke_26_rtheta_kernel_type' dm=True]
    HaloExchange[field='rtheta', type='region', depth=1, check_dirty=True]
    HaloExchange[field='theta', type='region', depth=1, check_dirty=True]
    HaloExchange[field='wind', type='region', depth=1, check_dirty=True]
    Loop[type='',field_space='w0',it_space='cells', upper_bound='cell_halo(1)']
        KernCall rtheta_code(rtheta,theta,wind) [module_inline=False]
```

Simple python script to apply Open MP transformation
Can apply on whole model
Or as fine-grained as single file

```
from psyclone.transformations import Dynamo0p3ColourTrans, \
Dynamo0p3OMPLoopTrans, \
OMPParallelTrans

def trans(psy):
    ctrans = Dynamo0p3ColourTrans()
    otrans = Dynamo0p3OMPLoopTrans()
    oregtrans = OMPParallelTrans()

    # Loop over all of the Invokes in the PSy object
    for invoke in psy.invokes.invoke_list:

        print "Transforming invoke '{0}' ...".format(invoke.name)
        schedule = invoke.schedule

        # Colour loops unless they are on W3 or over dofs
        for loop in schedule.loops():
            if loop.iteration_space == "cells" and loop.field_space != "w3":
                schedule, _ = ctrans.apply(loop)

        # Add OpenMP to loops unless they are over colours
        for loop in schedule.loops():
            if loop.loop_type != "colours":
                schedule, _ = oregtrans.apply(loop)
                schedule, _ = otrans.apply(loop, reprod=True)

        # take a look at what we've done
        schedule.view()

    return psy
```

```
C Transforming invoke 'invoke_26_rtheta_kernel_type' ...
Schedule[invoke='invoke_26_rtheta_kernel_type' dm=True]
    HaloExchange[field='rtheta', type='region', depth=1, check_dirty=True]
    HaloExchange[field='theta', type='region', depth=1, check_dirty=True]
    HaloExchange[field='wind', type='region', depth=1, check_dirty=True]
    Loop[type='colours',field_space='w0',it_space='cells', upper_bound='ncolours']
        Directive[OMP parallel]
        Directive[OMP do]
            Loop[type='colour',field_space='w0',it_space='cells', upper_bound='ncolour']
                KernCall rtheta_code(rtheta,theta,wind) [module_inline=False]
```


Update halos
YAXT→MPI

colouring from
infrastructure


OpenMP
workshare across
cells in colour

```
!
IF (chi_proxy(3)%is_dirty(depth=1)) THEN
  CALL chi_proxy(3)%halo_exchange(depth=1)
END IF
!
CALL rhs_heldsuarez_proxy%vspace%get_colours(ncolour, ncp_colour, cmap)
!
DO colour=1,ncolour
  !$omp parallel default(shared), private(cell)
  !$omp do schedule(static)
  DO cell=1,ncp_colour(colour)
    !
    call held_suarez_code(nlayers, rhs_heldsuarez_proxy%data, &
      ! ...
      )
  END DO
  !$omp end do
  !$omp end parallel
END DO
!
! Set halos dirty for fields modified in the above loop
!
CALL rhs_heldsuarez_proxy%set_dirty()
```


kernel call for single
column. Args are
arrays and scalars




Dof living on shared (partitioned) entity (edge).
Receive contribution from owned and halo cell.
Redundant compute contribution in halo to shared dof.
Less communication




MPI only, 4 MPI ranks all have halos
Hybrid, 1 MPI task has a halo, 4 OpenMP threads share halo boundary-to-area scaling
→ Less work for OpenMP threads





Point-wise computations (e.g. set field to a scalar) loop over dofs
Looping to owned dofs → halo exchange required for P2
Looping to annexed dofs is now transformation in Psyclone
Small increase in redundant computation
Large reduction in number of halo exchanges required



C576 is 576x576x6 cubed sphere ~ 17Km resolution
Intel 17, Cray XC40 dual socket, 18-core Broadwell
MPI is 36 MPI ranks per node

OMP is 6 MPI ranks / 6 omp threads


MV is matrix-vector kernel
HE is Halo-Exchange

Redundant computation

- i. 7.5 x reduction in number of halo exchanges
- ii. OMP has less work than MPI

Only compile and run code with Intel Compiler
Cray, PGI have problems with F2K3 OO code
Intel 17 OMP profile shows ~10% run-time is OMP synchronisation.
Grows with number of MPI ranks
`OMP_WAIT_POLICY=active` helps a bit
c.f. CCE8.5.8 and Intel 17 on single kernel code show 1-2% OMP synchronisation
Not clear what (and whose) the problem is?

Cray XC40 Aries network variability affects local comms (Halo exchange) and global comms (global sum)
Hard to measure performance without doing lots of runs
Is OMP faster than MPI. It can be!



C1152 cubed sphere ~ 9 Km
Comparing Multigrid and
Krylov subspace solvers
Scales well (out to 55K
cores) – mixed mode
LV at far right is 12x12 x 30L
For bigger problems can
scale to more nodes
Target is 1Km resolution
MS23: CMM: Thurs 1245 -
*Scalable Linear Solvers for
Next Generation Weather
and Climate Models*

Heterogeneous nodes with distinct memory spaces .e.g GPU
OpenACC and OpenMP4.5 (++) offload kernel for execution
PGI and Cray have problems compiling F2K3 OO code
PSyclone can generate OpenACC in PSy layer (GOcean)
Cannot yet annotate kernel code - ongoing

PSyclone Kernel Extractor (PSYKE) – dump out looping data in PSy layer
Dummy PSy layer (driver) can run a single kernel in isolation
LFRic-microbenchmark suite (On Github)
Can experiment with single kernel code

Offload data and kernel, same logic as OpenMP
Want bigger data regions


```
$acc loop vector private(ik, df, df2, lhs_e, x_e)
do k = 0, nlayers-1
  do df = 1, ndf2
    x_e(df) = x(map2(df)+k)
  end do
  lhs_e(:) = 0.0_r_def
  ik = (cell-1)*nlayers + k + 1
  do df = 1, ndf1
    do df2 = 1, ndf2
      lhs_e(df) = lhs_e(df) + matrix(df,df2,ik)*x_e(df2)
    end do
  end do
  do df = 1,ndf1
    !$acc atomic update
    lhs(map1(df)+k) = lhs(map1(df)+k) + lhs_e(df)
  end do
end do
$acc end loop
```

```
!$acc data copyin(ptheta_2_local_stencil,x_data,map_any_
!$acc& ncp_colour, nlayers,ncell_3d,ndf_any_space_1_theta_
!$acc& ndf_any_space_2_x, undf_any_space_2_x) copy(theta_
do colour = 1, ncolour
 !$acc parallel loop private(cell, map1, map2), firstprivat
  do cell = 1, ncp_colour(colour)

    map1(:)=map_any_space_1_theta_adv_term(:,cmap(colou
    map2(:)=map_any_space_2_x(:,cmap(colour,cell))
    call matrix_vector_code(cell, nlayers, theta_adv_t
      ptheta_2_local_stencil, ndf_any_space_1_theta_
      undf_any_space_1_theta_adv_term, &
      map1, &
      ndf_any_space_2_x, undf_any_space_2_x,&
      map2 )
    end do
    !$acc end parallel loop
  end do
 !$acc end data
```

need to annotate kernel source
SIMD (vector/warp) level
parallelism

A. Gray (NVIDIA)
Cumulative speed up
against original
OpenACC code.
Problem size is too
small for GPU.
Amortise cost of data
movement by
offloading multiple
kernels



Summary

Separation of concerns is a powerful abstraction

High-level language + Optimising compiler is not sufficient in the age of parallelism


Lots of programming models – no single one is sufficient

Developing DSL - PSyKAI API + PSyclone

Performance, Portability and Productivity



32 core per socket
Can over-subscribe
4 HDW threads
Cray Compiler
PE decrease due to
small problem size

Parallel efficiency c.f. no opt 1 thread



Data layout, unstructured mesh

W_0 space (vertices)



Data array (1-d)

Dofmap 2-d array



```
data (map (1 , 4) + 0) ←  
data (map (1 , 3) + 1) ←  
data (map (2 , 2) + 0) ←  
data (map (2 , 1) + 1) ←
```

n_{cell}

PSy layer

Kernel layer

Visit same dof more than once: loop over cells, levels, dofs
 Mesh and dofmap form an ordered set
 Change mesh topology (element), geometry (cubed sphere)
 → Change to mesh generation and partition
 → No change to science code